

The screenshot shows the AWS CloudWatch Metrics Insights interface. It has two main tabs: "Tasks" and "Services".

**Tasks Tab:**

- Last updated: February 13, 2026, 22:39 (UTC-8:00)
- Filter desired status: Any desired status
- Filter launch type: Any launch type
- Task details: Task ID 4d55842948f0453fa0db6..., Desired status Running, Task definition CS6650L2-task:1, Health status Unknown, Created at 3 minutes ago.

**Services Tab:**

- Last updated: February 13, 2026, 22:39 (UTC-8:00)
- Filter launch type: Any launch type
- Filter scheduling strategy: Any scheduling strategy
- Filter resource management type: Any resource management type
- Service details: Service name CS6650L2, ARN arn:aws:ecs:us-v, Status Active, Type REPLICA, Last deployment CS6650L2...

Http response:

# 204 — Success (add a product)

```
● (base) emilychen@emilys-Air-2 terraform % curl -v -X POST http://18.237.155.56:8080/products/1/details \
-H "Content-Type: application/json" \
-d '{"product_id":1,"sku":"ABC-123-XYZ","manufacturer":"Acme Corporation","category_id":456,"weight":1250,"some_other_id":789}' \
Note: Unnecessary use of -X or --request, POST is already inferred.
* Trying 18.237.155.56:8080...
* Connected to 18.237.155.56 (18.237.155.56) port 8080
> POST /products/1/details HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: /*
> Content-Type: application/json
> Content-Length: 122
>
< HTTP/1.1 204 No Content
< Date: Sat, 14 Feb 2026 03:27:00 GMT
<
* Connection #0 to host 18.237.155.56 left intact
```

# 400 — Missing required field

```
ome_other_id":789}'
```

- (base) emilychen@emilys-Air-2 terraform % curl -v -X POST http://18.237.155.56:8080/products/1/details \  
-H "Content-Type: application/json" \  
-d '{"product\_id":1,"sku":"ABC-123-XYZ"}'  
Note: Unnecessary use of -X or --request, POST is already inferred.  
\* Trying 18.237.155.56:8080...  
\* Connected to 18.237.155.56 (18.237.155.56) port 8080  
> POST /products/1/details HTTP/1.1  
> Host: 18.237.155.56:8080  
> User-Agent: curl/8.4.0  
> Accept: \*/\*  
> Content-Type: application/json  
> Content-Length: 36  
>  
< HTTP/1.1 400 Bad Request  
< Content-Type: application/json  
< Date: Sat, 14 Feb 2026 03:27:15 GMT  
< Content-Length: 63  
<  
{"error":"INVALID\_INPUT","message":"manufacturer is required"}

#### # 400 — Invalid JSON

```
curl -v -X POST http://18.237.155.56:8080/products/1/details \  
-H "Content-Type: application/json" \  
-d 'not json'  
zsh: command not found: #  
Note: Unnecessary use of -X or --request, POST is already inferred.  
* Trying 18.237.155.56:8080...  
* Connected to 18.237.155.56 (18.237.155.56) port 8080  
> POST /products/1/details HTTP/1.1  
> Host: 18.237.155.56:8080  
> User-Agent: curl/8.4.0  
> Accept: */*  
> Content-Type: application/json  
> Content-Length: 8  
>  
< HTTP/1.1 400 Bad Request  
< Content-Type: application/json  
< Date: Sat, 14 Feb 2026 03:27:40 GMT  
< Content-Length: 106  
<  
{"error":"INVALID_INPUT","message":"Invalid JSON: invalid character 'o' in literal null (expecting 'u')"}  
* Connection #0 to host 18.237.155.56 left intact
```

#### # 400 — Non-integer product ID

```
● (base) emilychen@emilys-Air-2 terraform % curl -v -X POST http://18.237.155.56:8080/products/abc/details \  
-H "Content-Type: application/json" \  
-d '{}'  
Note: Unnecessary use of -X or --request, POST is already inferred.  
* Trying 18.237.155.56:8080...  
* Connected to 18.237.155.56 (18.237.155.56) port 8080  
> POST /products/abc/details HTTP/1.1  
> Host: 18.237.155.56:8080  
> User-Agent: curl/8.4.0  
> Accept: */*  
> Content-Type: application/json  
> Content-Length: 2  
>  
< HTTP/1.1 400 Bad Request  
< Content-Type: application/json  
< Date: Sat, 14 Feb 2026 03:28:45 GMT  
< Content-Length: 67  
<  
{"error":"INVALID_INPUT","message":"productId must be an integer"}
```

#### # 400 — Product ID below minimum

```
● (base) emilychen@emilys-Air-2 terraform % curl -v -X POST http://18.237.155.56:8080/products/0/details \
  -H "Content-Type: application/json" \
  -d '{}'
Note: Unnecessary use of -X or --request, POST is already inferred.
*   Trying 18.237.155.56:8080...
*   Connected to 18.237.155.56 (18.237.155.56) port 8080
> POST /products/0/details HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: */*
> Content-Type: application/json
> Content-Length: 2
>
< HTTP/1.1 400 Bad Request
< Content-Type: application/json
< Date: Sat, 14 Feb 2026 03:28:50 GMT
< Content-Length: 66
<
{"error":"INVALID_INPUT","message":"productId must be \u003e= 1"}
* Connection #0 to host 18.237.155.56 left intact
```

#### # 200 — Product found

```
● (base) emilychen@emilys-Air-2 terraform % curl -v http://18.237.155.56:8080/products/1
*   Trying 18.237.155.56:8080...
*   Connected to 18.237.155.56 (18.237.155.56) port 8080
> GET /products/1 HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: */*
>
< HTTP/1.1 200 OK
< Content-Type: application/json
< Date: Sat, 14 Feb 2026 03:28:56 GMT
< Content-Length: 123
<
{"product_id":1,"sku":"ABC-123-XYZ","manufacturer":"Acme Corporation","category_id":456,"weight":1250,"some_other_id":789}
* Connection #0 to host 18.237.155.56 left intact
```

#### # 404 — Product not found

```
● (base) emilychen@emilys-Air-2 terraform % curl -v http://18.237.155.56:8080/products/999
*   Trying 18.237.155.56:8080...
*   Connected to 18.237.155.56 (18.237.155.56) port 8080
> GET /products/999 HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: */*
>
< HTTP/1.1 404 Not Found
< Content-Type: application/json
< Date: Sat, 14 Feb 2026 03:29:04 GMT
< Content-Length: 64
<
>{"error":"NOT_FOUND","message":"product with ID 999 not found"}
* Connection #0 to host 18.237.155.56 left intact
```

#### # 400 — Non-integer product ID

```
● (base) emilychen@emilys-Air-2 terraform % curl -v http://18.237.155.56:8080/products/abc
*   Trying 18.237.155.56:8080...
* Connected to 18.237.155.56 (18.237.155.56) port 8080
> GET /products/abc HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: */*
>
< HTTP/1.1 400 Bad Request
< Content-Type: application/json
< Date: Sat, 14 Feb 2026 03:29:12 GMT
< Content-Length: 67
<
{"error":"INVALID_INPUT","message":"productId must be an integer"}
* Connection #0 to host 18.237.155.56 left intact
```

# 400 — Product ID below minimum

```
● (base) emilychen@emilys-Air-2 terraform % curl -v http://18.237.155.56:8080/products/0
*   Trying 18.237.155.56:8080...
* Connected to 18.237.155.56 (18.237.155.56) port 8080
> GET /products/0 HTTP/1.1
> Host: 18.237.155.56:8080
> User-Agent: curl/8.4.0
> Accept: */*
>
< HTTP/1.1 400 Bad Request
< Content-Type: application/json
< Date: Sat, 14 Feb 2026 03:29:19 GMT
< Content-Length: 66
<
{"error":"INVALID_INPUT","message":"productId must be \u003e= 1"}
* Connection #0 to host 18.237.155.56 left intact
/home/emilychen/Downloads/Air-2_terraform %
```

Test experiments (60s)

10 user, 2 ramp up


HttpUser:

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS




Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/products/{id}	265	0	17	22	24	17.67	13	31	118.45	4.8	0
POST	/products/{id}/details	37	0	17	26	31	18.07	14	31	0	0.5	0
POST	/products/{id}/details (seed)	1000	0	16	21	37	16.81	12	54	0	0	0
Aggregated		1302	0	16	21	35	17.02	12	54	24.11	5.3	0


### Total Requests per Second



### Response Times (ms)



### Number of Users



FastHttpUser:

**LOCUST**

Host: http://18.237.155.56:8080 Status: STOPPED RPS: 5.1 Failures: 0% NEW RESET

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/products/{id}	258	0	16	19	20	15.93	12	20	118.45	4.8	0
POST	/products/{id}/details	29	0	16	19	26	16.31	13	26	0	0.3	0
POST	/products/{id}/details (seed)	1000	0	15	18	30	15.19	12	36	0	0	0
Aggregated		1287	0	15	18	21	15.36	12	36	23.74	5.1	0


**Total Requests per Second**

**Response Times (ms)**

**Number of Users**

100 user, 10 ramp up

## HttpUser:




## FastHttpUser:


**LOCUST**

Host http://18.237.155.56:8080										Status STOPPED	RPS 51	Failures 0%	<b>NEW</b>	<b>RESET</b>	
<b>STATISTICS</b>		CHARTS	FAILURES	EXCEPTIONS	CURRENT RATIO	DOWNLOAD DATA			LOGS						
Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s			
GET	/products/{id}	2424	0	15	18	24	15.23	12	141	118.47	44.8	0			
POST	/products/{id}/details	290	0	15	18	24	15.25	12	54	0	6.2	0			
POST	/products/{id}/details (seed)	10000	0	16	25	44	17.02	11	149	0	0	0			
Aggregated		12714	0	15	23	42	16.64	11	149	22.59	51	0			


**Total Requests per Second**



**Response Times (ms)**



**Number of Users**



500 user, 50 ramp up  
HttpUser



**LOCUST**

Host: http://18.237.155.56:8080 Status: STOPPED RPS: 247.8 Failures: 0% NEW RESET

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS

III

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/products/{id}	13479	0	17	48	84	20.1	12	204	118.37	222.2	0
POST	/products/{id}/details	1501	0	17	46	85	20.24	12	113	0	25.6	0
POST	/products/{id}/details (seed)	50000	0	63	110	140	65.65	13	378	0	0	0
Aggregated		64980	0	53	110	130	55.15	12	378	24.55	247.8	0



FastHttpUser:

**LOCUST**

Host: http://18.237.155.56:8080 Status: STOPPED RPS: 247.9 Failures: 0% NEW RESET ⚙️

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/products/{id}	11446	0	15	48	89	18.6	11	175	118.48	222.8	0
POST	/products/{id}/details	1263	0	15	58	88	18.93	11	138	0	25.2	0
POST	/products/{id}/details (seed)	50000	0	69	100	110	56.44	12	181	0	0	0
Aggregated		62709	0	24	100	110	48.78	11	181	21.63	247.9	0

**Total Requests per Second**

**Response Times (ms)**

**Number of Users**

1000 user, 100 ramp up  
HttpUser:

**LOCUST**

Host: http://18.237.155.56:8080 Status: STOPPED RPS: 492.7 Failures: 0% NEW RESET ⚙️

STATISTICS CHARTS FAILURES EXCEPTIONS CURRENT RATIO DOWNLOAD DATA LOGS

Type	Name	# Requests	# Fails	Median (ms)	95%ile (ms)	99%ile (ms)	Average (ms)	Min (ms)	Max (ms)	Average size (bytes)	Current RPS	Current Failures/s
GET	/products/{id}	16940	0	17	180	240	37.38	12	336	118.33	440.3	0
POST	/products/{id}/details	1925	0	17	190	240	37.5	13	349	0	52.4	0
POST	/products/{id}/details (seed)	100000	0	200	290	340	183.54	13	1142	0	0	0
Aggregated		118865	0	180	280	340	160.34	12	1142	16.86	492.7	0

**Total Requests per Second**

RPS Failures/s

**Response Times (ms)**

50th percentile 95th percentile

**Number of Users**


Number of Users

```

(base) emilys-Air-2:~ emilys % locust -f locustfile.py --host http://18.237.155.56:8080 --httpUserLoadTest
[2026-02-13 22:21:10,267] emilys-Air-2/INFO/locust.main: Starting Locust 2.43.3
[2026-02-13 22:21:10,279] emilys-Air-2/INFO/locust.main: Starting web interface at http://0.0.0.0:8089, press
enter to open your default browser.
[2026-02-13 22:21:38,697] emilys-Air-2/INFO/locust.runners: Ramping to 1000 users at a rate of 100.00 per seco
nd
[2026-02-13 22:21:48,015] emilys-Air-2/INFO/locust.runners: All users spawned: {"HttpUserLoadTest": 1000} (100
0 total users)
[2026-02-13 22:21:50,425] emilys-Air-2/WARNING/root: CPU usage above 90%! This may constrain your throughput a
nd may even give inconsistent response time measurements! See https://docs.locust.io/en/stable/running-distribu
ted.html for how to distribute the load over multiple CPU cores or machines
[2026-02-13 22:22:39,276] emilys-Air-2/WARNING/locust.runners: CPU usage was too high at some point during the
test! See https://docs.locust.io/en/stable/running-distributed.html for how to distribute the load over multi
ple CPU cores or machines

```

FastHttpUser:



## Summary Table

User	Type	GET Median	GET p95	POST Median	POST p95	RPS	Failures/s
------	------	------------	---------	-------------	----------	-----	------------

10	Http	17ms	22ms	17ms	26ms	5.3	0%
10	Fast	16ms	19ms	16ms	19ms	5.1	0%
100	Http	17ms	22ms	17ms	21ms	48.6	0%
100	Fast	15ms	18ms	15ms	18ms	51	0%
500	Http	15ms	48ms	15ms	58ms	247.8	0%
500	Fast	17ms	48ms	17ms	48ms	247.9	0%
1000	Http	17ms	180ms	17ms	190ms	492.7	0%
1000	Fast	15ms	120ms	15ms	140ms	497.3	0%

## Key Findings

- 1. The server never broke** : 0% failures across all tests, even at 1000 users. The Go server with the in-memory hashmap is efficient. This makes sense: Go's goroutine-per-request model handles concurrency well, and the `RWMutex` + hashmap has virtually zero contention.
- 2. Median response times stayed flat (~15-17ms)** : this is basically the network round-trip time from local machine to AWS. The server itself responds nearly instantly; what we're measuring is mostly wire latency. This is a sign of a healthy server that isn't saturating.
- 3. Tail latencies (p95/p99) varies**: while medians stayed flat, the p95 jumped dramatically at high concurrency: 22ms at 10 users → 180ms at 1000 users for HttpUser. This means most requests are fast, but some get queued and wait longer under load.
- 4. HttpUser vs FastHttpUser — the difference appears at 1000 users:**
  - p95 latency: **180ms (Http) vs 120ms (Fast)** — FastHttpUser is 33% faster at the tail
  - p99 latency: **240ms (Http) vs 240ms (Fast)** — converges at the extreme tail
  - The HttpUser 1000-user test even threw a **CPU usage warning** ("CPU usage above 90%"), while FastHttpUser did not. This is the key insight: at high concurrency, HttpUser's Python `requests` library consumes more CPU on the local machine, which inflates response times. FastHttpUser's C-optimized parser handles the same load with less overhead.
- 5. Why no difference at low concurrency?** At 10-100 users, both clients are mostly idle between requests (due to `wait_time`). The bottleneck is network latency (~15ms), not local CPU. FastHttpUser's efficiency only matters when the local machine is generating hundreds of requests per second.

**6. GET vs POST performance is nearly identical** — both have the same median. This is because the in-memory hashmap operations (read or write) are sub-microsecond. The [RWMutex](#) benefit would become visible with a slower data store (like a database) where read locks allow true parallelism while write locks block.

**7. RPS scaled linearly** — 5 → 49 → 248 → 497 RPS as users increased 10 → 100 → 500 → 1000. The server hasn't hit its ceiling yet. We'd likely need thousands more users to find the breaking point.